

Elearning data analysis

2025-02-14

Load packages

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.2.1
```

```
## — Attaching packages — tidyverse 1.3.2 —
## ✓ ggplot2 3.3.6      ✓ purrr   1.0.2
## ✓ tibble  3.1.8      ✓ dplyr   1.0.9
## ✓ tidyr   1.2.0      ✓ stringr 1.4.0
## ✓ readr   2.1.2      ✓ forcats 0.5.1
```

```
## Warning: package 'ggplot2' was built under R version 4.2.1
```

```
## Warning: package 'tibble' was built under R version 4.2.1
```

```
## Warning: package 'tidyr' was built under R version 4.2.1
```

```
## Warning: package 'readr' was built under R version 4.2.1
```

```
## Warning: package 'purrr' was built under R version 4.2.3
```

```
## Warning: package 'dplyr' was built under R version 4.2.1
```

```
## Warning: package 'forcats' was built under R version 4.2.1
```

```
## — Conflicts — tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
```

```
library(caret) #for classification and regression training
```

```
## Warning: package 'caret' was built under R version 4.2.1
```

```
## Loading required package: lattice
```

```
## Warning: package 'lattice' was built under R version 4.2.1
```

```
##  
## Attaching package: 'caret'  
##  
## The following object is masked from 'package:purrr':  
##  
## lift
```

```
library(ranger) #for random forests
```

```
## Warning: package 'ranger' was built under R version 4.2.3
```

```
library(e1071) #for statistics functions
```

```
## Warning: package 'e1071' was built under R version 4.2.1
```

```
library(tidylog)
```

```
##  
## Attaching package: 'tidylog'  
##  
## The following objects are masked from 'package:dplyr':  
##  
## add_count, add_tally, anti_join, count, distinct, distinct_all,  
## distinct_at, distinct_if, filter, filter_all, filter_at, filter_if,  
## full_join, group_by, group_by_all, group_by_at, group_by_if,  
## inner_join, left_join, mutate, mutate_all, mutate_at, mutate_if,  
## relocate, rename, rename_all, rename_at, rename_if, rename_with,  
## right_join, sample_frac, sample_n, select, select_all, select_at,  
## select_if, semi_join, slice, slice_head, slice_max, slice_min,  
## slice_sample, slice_tail, summarise, summarise_all, summarise_at,  
## summarise_if, summarize, summarize_all, summarize_at, summarize_if,  
## tally, top_frac, top_n, transmute, transmute_all, transmute_at,  
## transmute_if, ungroup  
##  
## The following objects are masked from 'package:tidyr':  
##  
## drop_na, fill, gather, pivot_longer, pivot_wider, replace_na,  
## spread, uncount  
##  
## The following object is masked from 'package:stats':  
##  
## filter
```

```
library(dataedu) #package for educational data
```

Import and view data

```
setwd("D:/Data-analytics-project/Elearning-data-analysis")  
df <- dataedu::sci_mo_with_text  
glimpse(df)
```

```

## Rows: 606
## Columns: 74
## $ student_id      <dbl> 43146, 44638, 47448, 47979, 48797, 51943, 52326,...
## $ course_id       <chr> "FrScA-S216-02", "OcnA-S116-01", "FrScA-S216-01"...
## $ total_points_possible <dbl> 3280, 3531, 2870, 4562, 2207, 4208, 4325, 2086, ...
## $ total_points_earned <dbl> 2220, 2672, 1897, 3090, 1910, 3596, 2255, 1719, ...
## $ percentage_earned <dbl> 0.6768293, 0.7567261, 0.6609756, 0.6773345, 0.86...
## $ subject         <chr> "FrScA", "OcnA", "FrScA", "OcnA", "PhysA", "FrSc...
## $ semester        <chr> "S216", "S116", "S216", "S216", "S116", "S216", ...
## $ section         <chr> "02", "01", "01", "01", "01", "03", "01", "01", ...
## $ Gradebook_Item  <chr> "POINTS EARNED & TOTAL COURSE POINTS", "ATTEMPT...
## $ Grade_Category  <lg1> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
## $ final_grade     <dbl> 93.45372, 81.70184, 88.48758, 81.85260, 84.00000...
## $ Points_Possible <dbl> 5, 10, 10, 5, 438, 5, 10, 10, 443, 5, 12, 10, 5,...
## $ Points_Earned   <dbl> NA, 10.00, NA, 4.00, 399.00, NA, NA, 10.00, 425...
## $ Gender          <chr> "M", "F", "M", "M", "F", "F", "M", "F", "F", "M"...
## $ q1              <dbl> 5, 4, 5, 5, 4, NA, 5, 3, 4, NA, NA, 4, 3, 5, NA,...
## $ q2              <dbl> 4, 4, 4, 5, 3, NA, 5, 3, 3, NA, NA, 5, 3, 3, NA,...
## $ q3              <dbl> 4, 3, 4, 3, 3, NA, 3, 3, 3, NA, NA, 3, 3, 5, NA,...
## $ q4              <dbl> 5, 4, 5, 5, 4, NA, 5, 3, 4, NA, NA, 5, 3, 5, NA,...
## $ q5              <dbl> 5, 4, 5, 5, 4, NA, 5, 3, 4, NA, NA, 5, 4, 5, NA,...
## $ q6              <dbl> 5, 4, 4, 5, 4, NA, 5, 4, 3, NA, NA, 5, 3, 5, NA,...
## $ q7              <dbl> 5, 4, 4, 4, 4, NA, 4, 3, 3, NA, NA, 5, 3, 5, NA,...
## $ q8              <dbl> 5, 5, 5, 5, 4, NA, 5, 3, 4, NA, NA, 4, 3, 5, NA,...
## $ q9              <dbl> 4, 4, 3, 5, NA, NA, 5, 3, 2, NA, NA, 5, 2, 2, NA...
## $ q10             <dbl> 5, 4, 5, 5, 3, NA, 5, 3, 5, NA, NA, 4, 4, 5, NA,...
## $ time_spent       <dbl> 1555.1667, 1382.7001, 860.4335, 1598.6166, 1481....
## $ TimeSpent_hours <dbl> 25.91944500, 23.04500167, 14.34055833, 26.643610...
## $ TimeSpent_std    <dbl> -0.18051496, -0.30780313, -0.69325954, -0.148446...
## $ int             <dbl> 5.0, 4.2, 5.0, 5.0, 3.8, 4.6, 5.0, 3.0, 4.2, NA,...
## $ pc              <dbl> 4.50, 3.50, 4.00, 3.50, 3.50, 4.00, 3.50, 3.00, ...
## $ uv              <dbl> 4.333333, 4.000000, 3.666667, 5.000000, 3.500000...
## $ enrollment_status <chr> "Approved/Enrolled", "Approved/Enrolled", "Appro...
## $ enrollment_reason <chr> "Course Unavailable at Local School", "Course Un...
## $ cogproc         <dbl> 15.069737, 7.106667, 15.165854, 14.508000, 16.69...
## $ male            <dbl> 0.51210526, 0.00000000, 0.11121951, 0.00000000, ...
## $ female          <dbl> 0.16657895, 0.00000000, 0.15219512, 0.00000000, ...
## $ friend          <dbl> 0.00000000, 0.00000000, 0.01268293, 0.00000000, ...
## $ family          <dbl> 0.006052632, 0.00000000, 0.084878049, 0.00000000...
## $ social          <dbl> 6.200526, 6.140000, 5.052927, 6.133000, 7.534000...
## $ sad             <dbl> 0.18078947, 0.00000000, 0.09097561, 0.00000000, ...
## $ anger           <dbl> 0.41868421, 0.00000000, 0.14097561, 0.10800000, ...
## $ anx             <dbl> 0.08000000, 0.00000000, 0.275365854, 0.7880000...
## $ negemo          <dbl> 1.1363158, 0.0000000, 1.4187805, 1.1520000, 1.28...
## $ posemo          <dbl> 3.555526, 19.010000, 2.906098, 5.591000, 3.79400...
## $ affect          <dbl> 4.756053, 19.010000, 4.330732, 6.743000, 5.07500...
## $ quant           <dbl> 2.046842, 2.743333, 3.245366, 3.214000, 2.551000...
## $ number          <dbl> 0.9131579, 3.4733333, 2.3065854, 0.2570000, 0.21...
## $ interrog        <dbl> 1.2857895, 0.4433333, 1.7868293, 1.1030000, 1.71...
## $ compare         <dbl> 2.4213158, 4.1466667, 3.9021951, 2.6990000, 3.94...
## $ adj             <dbl> 5.106842, 5.480000, 5.614390, 5.213000, 4.618000...
## $ verb            <dbl> 18.11368, 11.02333, 16.34366, 16.31100, 17.11700...

```

```
## $ negate <dbl> 1.2060526, 0.0000000, 1.6809756, 1.1300000, 0.74...
## $ conj <dbl> 5.565526, 6.660000, 5.370244, 6.203000, 7.244000...
## $ adverb <dbl> 6.243421, 6.660000, 5.824878, 5.314000, 6.492000...
## $ auxverb <dbl> 11.298421, 9.246667, 10.226341, 8.890000, 9.4940...
## $ prep <dbl> 12.301579, 11.850000, 12.132927, 13.626000, 12.8...
## $ article <dbl> 7.828947, 2.223333, 6.767805, 9.119000, 9.830000...
## $ ipron <dbl> 6.936316, 2.743333, 5.145122, 4.335000, 7.841000...
## $ they <dbl> 1.01026316, 0.00000000, 0.84341463, 1.86300000, ...
## $ shehe <dbl> 0.54342105, 0.00000000, 0.16951220, 0.00000000, ...
## $ you <dbl> 1.7442105, 3.4733333, 1.1487805, 2.0490000, 2.62...
## $ we <dbl> 0.06578947, 0.00000000, 0.03317073, 0.30200000, ...
## $ i <dbl> 3.646579, 7.993333, 4.689268, 3.449000, 3.142000...
## $ ppron <dbl> 7.010000, 11.470000, 6.882927, 7.662000, 6.77900...
## $ pronoun <dbl> 13.98868, 14.20667, 12.02756, 12.21900, 14.61900...
## $ `function` <dbl> 55.15447, 44.63000, 49.40293, 53.12700, 57.50900...
## $ Dic <dbl> 86.27895, 86.31000, 80.72220, 86.49700, 90.48700...
## $ Sixltr <dbl> 20.89316, 22.20333, 20.80780, 21.80200, 15.30600...
## $ WPS <dbl> 17.413947, 9.833333, 17.922439, 18.824000, 15.66...
## $ Tone <dbl> 56.62395, 96.38000, 49.41610, 78.36900, 55.38400...
## $ Authentic <dbl> 44.13079, 70.25333, 41.22366, 49.03800, 42.25000...
## $ Clout <dbl> 49.52079, 53.58333, 40.11024, 53.08800, 54.08500...
## $ Analytic <dbl> 55.70316, 56.04000, 58.98098, 69.95700, 55.82000...
## $ WC <dbl> 88.31579, 34.66667, 69.34146, 61.20000, 47.10000...
## $ n <dbl> 38, 3, 41, 10, 10, 2, 21, 18, 31, 37, 37, 18, 12...
```

Data processing

Select important variables

```
df <-
df %>%
select(
  int, #student think this course is interesting
  uv, #utility value: what I'm learning in this course is relevant to my life
  pc, #perceived competence: this topic is one of my best subjects
  time_spent, #time spent in the course
  final_grade,
  subject,
  enrollment_reason,
  semester,
  enrollment_status,
  cogproc, #student's cognitive processing
  social, #social-related discourse in discussion board posts
  posemo, #positive emotions in discussion board posts
  negemo, #negative emotions in discussion board posts
  n #number of discussion board posts in the course in the semester
)
```

```
## select: dropped 60 variables (student_id, course_id, total_points_possible,
## total_points_earned, percentage_earned, ...)
```

Analysis

Analyze data

```
# Check number of rows in dataset  
nrow(df)
```

```
## [1] 606
```

```
# Drop rows with missing data (N/A)  
df<-na.omit(df)  
# Check number of rows after dropping N/A  
nrow(df)
```

```
## [1] 464
```

```
glimpse(df)
```

```
## Rows: 464  
## Columns: 14  
## $ int          <dbl> 5.0, 4.2, 5.0, 5.0, 3.8, 5.0, 3.0, 4.2, 4.4, 3.4, 4.0...  
## $ uv           <dbl> 4.333333, 4.000000, 3.666667, 5.000000, 3.500000, 5.0...  
## $ pc           <dbl> 4.50, 3.50, 4.00, 3.50, 3.50, 3.50, 3.00, 3.00, 4.00...  
## $ time_spent   <dbl> 1555.1667, 1382.7001, 860.4335, 1598.6166, 1481.8000...  
## $ final_grade  <dbl> 93.45372, 81.70184, 88.48758, 81.85260, 84.00000, 83...  
## $ subject      <chr> "FrScA", "OcnA", "FrScA", "OcnA", "PhysA", "AnPhA", ...  
## $ enrollment_reason <chr> "Course Unavailable at Local School", "Course Unavai...  
## $ semester     <chr> "S216", "S116", "S216", "S216", "S116", "S216", "S11...  
## $ enrollment_status <chr> "Approved/Enrolled", "Approved/Enrolled", "Approved/...  
## $ cogproc      <dbl> 15.069737, 7.106667, 15.165854, 14.508000, 16.692000...  
## $ social       <dbl> 6.200526, 6.140000, 5.052927, 6.133000, 7.534000, 7.0...  
## $ posemo       <dbl> 3.555526, 19.010000, 2.906098, 5.591000, 3.794000, 5.0...  
## $ negemo       <dbl> 1.1363158, 0.0000000, 1.4187805, 1.1520000, 1.282000...  
## $ n            <dbl> 38, 3, 41, 10, 10, 21, 18, 31, 18, 12, 3, 16, 7, 42,...
```

```
# Determine if there are variables with no variability  
nearZeroVar(df, saveMetrics=TRUE)
```

##	freqRatio	percentUnique	zeroVar	nzv
## int	1.314815	9.0517241	FALSE	FALSE
## uv	1.533333	6.4655172	FALSE	FALSE
## pc	1.488372	3.8793103	FALSE	FALSE
## time_spent	1.000000	100.0000000	FALSE	FALSE
## final_grade	1.333333	93.1034483	FALSE	FALSE
## subject	1.648649	1.0775862	FALSE	FALSE
## enrollment_reason	3.154762	1.0775862	FALSE	FALSE
## semester	1.226601	0.6465517	FALSE	FALSE
## enrollment_status	0.000000	0.2155172	TRUE	TRUE
## cogproc	1.000000	96.9827586	FALSE	FALSE
## social	1.500000	96.1206897	FALSE	FALSE
## posemo	1.000000	96.7672414	FALSE	FALSE
## negemo	13.000000	90.7327586	FALSE	FALSE
## n	1.333333	10.1293103	FALSE	FALSE

zeroVar column of enrollment_status is True, so we will remove it. Variables with no variability may cause problems in some models.

```
df <-
  df %>%
  select(-enrollment_status)
```

```
## select: dropped one variable (enrollment_status)
```

```
# Convert string categorical variables to factors
df <-
  df %>%
  mutate_if(is.character, as.factor)
```

```
## mutate_if: converted 'subject' from character to factor (0 new NA)
##           converted 'enrollment_reason' from character to factor (0 new NA)
##           converted 'semester' from character to factor (0 new NA)
```

Prepare train and test sets

```
# Set seed
set.seed(2025)
# Train 70%, Test 30%
# Create train set
trainIdx <- createDataPartition(df$final_grade,
                                p=.7,
                                list=FALSE,
                                times=1)

# Add new variable to dataset temporarily
# Select rows according to their row number
df <-
  df %>%
  mutate(temp_id = 1:464)
```

```
## mutate: new variable 'temp_id' (integer) with 464 unique values and 0% NA
```

```
# Filter dataset to get only rows indicated trainIdx vector
df_train <-
  df %>%
  filter(temp_id %in% trainIdx)
```

```
## filter: removed 136 rows (29%), 328 rows remaining
```

```
# Filter out to get test set
df_test <-
  df %>%
  filter(!temp_id %in% trainIdx)
```

```
## filter: removed 328 rows (71%), 136 rows remaining
```

```
# Delete temp_id from the original data
df <-
  df %>%
  select(-temp_id)
```

```
## select: dropped one variable (temp_id)
```

```
df_train <-
  df_train %>%
  select(-temp_id)
```

```
## select: dropped one variable (temp_id)
```



```
df_test <-  
  df_test %>%  
  select(-temp_id)
```

```
## select: dropped one variable (temp_id)
```

Estimate the model

Random forests - bootstrap resampling

```
# Set seed  
set.seed(2025)  
# Run random forest model  
# Final grade is y, other variables are x  
# Syntax: final_grade ~. => y ~ all other variables except y  
# Resampling method: Bootstrap resampling  
rf_fit <- train(final_grade ~.,  
                data = df_train,  
                method = "ranger")  
  
rf_fit
```

```
## Random Forest  
##  
## 328 samples  
## 12 predictor  
##  
## No pre-processing  
## Resampling: Bootstrapped (25 reps)  
## Summary of sample sizes: 328, 328, 328, 328, 328, 328, ...  
## Resampling results across tuning parameters:  
##  
##   mtry  splitrule  RMSE      Rsquared  MAE  
##    2    variance  16.07891  0.4814206  12.00116  
##    2  extratrees  17.37091  0.4558964  12.58493  
##   10    variance  14.81596  0.5150913  10.98432  
##   10  extratrees  14.37030  0.5813905  10.70215  
##   19    variance  14.88024  0.5093763  10.89464  
##   19  extratrees  13.78887  0.5957179  10.30089  
##  
## Tuning parameter 'min.node.size' was held constant at a value of 5  
## RMSE was used to select the optimal model using the smallest value.  
## The final values used for the model were mtry = 19, splitrule = extratrees  
## and min.node.size = 5.
```

```
# Results: Best RMSE: 13.79, Best R^2: 0.6
```

Random forests - Cross validation

```

#Set seed
set.seed(2025)
# Use cross validation
train_control <-
  trainControl(method="repeatedcv",
               number = 10, #10 folds
               repeats = 10) #repeat 10 times

rf_fit1 <-
  train(final_grade ~ .,
        data=df_train,
        method="ranger",
        trControl=train_control)

rf_fit1

```

```

## Random Forest
##
## 328 samples
## 12 predictor
##
## No pre-processing
## Resampling: Cross-Validated (10 fold, repeated 10 times)
## Summary of sample sizes: 294, 296, 294, 296, 295, 294, ...
## Resampling results across tuning parameters:
##
##  mtry  splitrule  RMSE      Rsquared  MAE
##    2    variance  15.48641  0.5306322  11.72206
##    2  extratrees  17.03196  0.5068189  12.51838
##   10    variance  13.90578  0.5720739  10.43246
##   10  extratrees  13.79550  0.6149065  10.45750
##   19    variance  13.80944  0.5732509  10.24774
##   19  extratrees  13.20623  0.6280336  10.01536
##
## Tuning parameter 'min.node.size' was held constant at a value of 5
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were mtry = 19, splitrule = extratrees
## and min.node.size = 5.

```

```

# Result: Best RMSE: 13.2, Best R^2: 0.63

```

Tuning random forest model

Previously: min.node.size is fixed to 5 Now: change min.node.size and mtry

```
set.seed(2025)

# Create a grid of different values of mtry, split rules and min node sizes to test
tune_grid <-
  expand.grid(
    mtry = c(2, 3, 7, 10, 19),
    splitrule = c("variance", "extratrees"),
    min.node.size=c(1, 5, 10, 15, 20)
  )

# Fit a new model using tuning grid
rf_fit2 <-
  train(final_grade~.,
    data=df_train,
    method="ranger",
    tuneGrid=tune_grid)

rf_fit2
```

```

## Random Forest
##
## 328 samples
## 12 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 328, 328, 328, 328, 328, 328, ...
## Resampling results across tuning parameters:
##
##      mtry  splitrule  min.node.size  RMSE      Rsquared  MAE
##      2    variance    1             15.99371  0.4861212  11.92869
##      2    variance    5             16.09348  0.4811132  12.01389
##      2    variance   10             16.18863  0.4746928  12.10158
##      2    variance   15             16.35733  0.4644302  12.23053
##      2    variance   20             16.46616  0.4581469  12.32386
##      2    extratrees  1             17.25220  0.4577243  12.47415
##      2    extratrees  5             17.37693  0.4552814  12.58400
##      2    extratrees 10             17.60935  0.4414375  12.76655
##      2    extratrees 15             17.78587  0.4314077  12.91368
##      2    extratrees 20             17.98658  0.4163722  13.07164
##      3    variance    1             15.53336  0.4947565  11.61643
##      3    variance    5             15.57180  0.4947766  11.64949
##      3    variance   10             15.69477  0.4856603  11.78690
##      3    variance   15             15.81486  0.4802026  11.89772
##      3    variance   20             15.99589  0.4696509  12.04563
##      3    extratrees  1             16.18111  0.5086335  11.73106
##      3    extratrees  5             16.35841  0.5011881  11.88576
##      3    extratrees 10             16.64786  0.4857649  12.12168
##      3    extratrees 15             16.90480  0.4740047  12.29219
##      3    extratrees 20             17.08276  0.4681437  12.43796
##      7    variance    1             14.91422  0.5138677  11.12181
##      7    variance    5             14.97616  0.5101048  11.16496
##      7    variance   10             15.02833  0.5065927  11.23460
##      7    variance   15             15.11661  0.5022217  11.34332
##      7    variance   20             15.20459  0.4968687  11.43802
##      7    extratrees  1             14.82990  0.5616703  10.94610
##      7    extratrees  5             14.87008  0.5631145  11.01215
##      7    extratrees 10             15.08067  0.5566656  11.17367
##      7    extratrees 15             15.33269  0.5447313  11.34670
##      7    extratrees 20             15.54238  0.5378536  11.49999
##     10    variance    1             14.77988  0.5181068  10.96699
##     10    variance    5             14.80325  0.5158682  10.99447
##     10    variance   10             14.81722  0.5147974  11.04178
##     10    variance   15             14.86410  0.5126755  11.09651
##     10    variance   20             14.93172  0.5092059  11.19363
##     10    extratrees  1             14.36338  0.5781267  10.67644
##     10    extratrees  5             14.39487  0.5797341  10.72813
##     10    extratrees 10             14.57807  0.5735542  10.86363
##     10    extratrees 15             14.78028  0.5663489  11.01265
##     10    extratrees 20             14.94316  0.5593999  11.14028
##     19    variance    1             14.88535  0.5095078  10.89231

```

```
## 19 variance 5 14.87368 0.5100980 10.89955
## 19 variance 10 14.88364 0.5086409 10.93401
## 19 variance 15 14.83982 0.5107337 10.91608
## 19 variance 20 14.79437 0.5138788 10.92149
## 19 extratrees 1 13.81893 0.5930555 10.32942
## 19 extratrees 5 13.82810 0.5940922 10.32892
## 19 extratrees 10 13.90730 0.5912578 10.40617
## 19 extratrees 15 13.99420 0.5891199 10.49275
## 19 extratrees 20 14.12019 0.5860253 10.60280
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were mtry = 19, splitrule = extratrees
## and min.node.size = 1.
```

```
# Result: Best MRSE: 13.82, Best R^2: 0.59
```

```
# See details of final model output of rf_fit2
rf_fit2$finalModel
```

```
## Ranger result
##
## Call:
## ranger::ranger(dependent.variable.name = ".outcome", data = x, mtry = min(param$mtry, n
col(x)), min.node.size = param$min.node.size, splitrule = as.character(param$splitrule), wr
ite.forest = TRUE, probability = classProbs, ...)
##
## Type: Regression
## Number of trees: 500
## Sample size: 328
## Number of independent variables: 19
## Mtry: 19
## Target node size: 1
## Variable importance mode: none
## Splitrule: extratrees
## Number of random splits: 1
## OOB prediction error (MSE): 174.9095
## R squared (OOB): 0.6188466
```

Examine predictive accuracy on test

```
set.seed(2025)
# Create new testing data including predicted values
df_test_augmented <-
  df_test %>%
  mutate(pred=predict(rf_fit2, df_test),
         obs=final_grade)
```

```
## mutate: new variable 'pred' (double) with 136 unique values and 0% NA
##          new variable 'obs' (double) with 133 unique values and 0% NA
```

```
# Transform object to data frame
defaultSummary(as.data.frame(df_test_augmented))
```

```
##          RMSE   Rsquared      MAE
## 12.0703944  0.6792062  9.0596247
```

- RMSE on test set = 12.07 => better than RMSE 13.82 of train set.
- R^2 on test set = 0.68 => better than R^2 0.59 of train set
- Therefore, the model performs well on unseen data (Test data)

Results

Variable Importance


```
set.seed(2025)
# Learn which variables contribute most strongly to the model prediction
rf_fit2_imp <-
  train(
    final_grade ~.,
    data=df_train,
    method="ranger",
    tuneGrid=tune_grid,
    importance="permutation"
  )

#Extract variable importance from the new model
varImp(rf_fit2_imp)
```

```
## ranger variable importance
##
## Overall
## n 100.0000
## subjectFrScA 25.9573
## time_spent 12.6614
## subjectPhysA 4.0058
## pc 4.0008
## negemo 3.6652
## semesterS216 3.2900
## social 1.6859
## enrollment_reasonOther 1.1195
## cogproc 1.0923
## subjectOcnA 0.9272
## enrollment_reasonScheduling Conflict 0.9037
## int 0.8653
## posemo 0.7464
## subjectBioA 0.7180
## enrollment_reasonLearning Preference of the Student 0.5989
## enrollment_reasonCredit Recovery 0.4383
## semesterT116 0.3946
## uv 0.0000
```

Visualize variable importance

```
varImp(rf_fit2_imp) %>%
  pluck(1) %>%
  rownames_to_column("var") %>%
  ggplot(aes(x = reorder(var, Overall), y = Overall)) +
  geom_col(fill = "#D55E00") +
  coord_flip() +
  theme_light()
```



Insights * Most important: n: number of student's discussion posts * 2nd most: subject Forensic Science: enrolled in Forensic Science course has great impact on student's * 3rd most: timespent: time student spent in the course

Compare random forest to regression

```
# Convert character variables to factors
df_train_lm <-
  df_train %>%
  mutate_if(is.character, as.factor)
```

```
## mutate_if: no changes
```

```
# Create a Linear regression model
lm_fit <-
  train(final_grade~.,
    data=df_train_lm,
    method="lm")
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```

```
## Warning in predict.lm(modelFit, newdata): prediction from a rank-deficient fit
## may be misleading
```



```
# Append predicted values to train set for linear model
df_train_lm <-
  df_train %>%
  mutate(obs=final_grade,
         pred=predict(lm_fit, df_train_lm))
```

```
## mutate: new variable 'obs' (double) with 312 unique values and 0% NA
##          new variable 'pred' (double) with 328 unique values and 0% NA
```

```
#Append predicted values to train set for random forest
df_train_rf <-
  df_train %>%
  mutate(pred=predict(rf_fit2, df_train),
         obs=final_grade)
```

```
## mutate: new variable 'pred' (double) with 328 unique values and 0% NA
##          new variable 'obs' (double) with 312 unique values and 0% NA
```

```
# Summarize Linear model
defaultSummary(as.data.frame(df_train_lm))
```

```
##          RMSE   Rsquared      MAE
## 15.0208404  0.5068252 11.2760167
```

```
# Summarize random forest
defaultSummary(as.data.frame(df_train_rf))
```

```
##          RMSE  Rsquared      MAE
## 4.8600844 0.9663748 3.6447460
```

Conclusions

- Random forest has higher R squared (0.97 compared to 0.51 for the regression model)
- Random forest has lower RMSE, meaning random forest fits the data better than linear model.